


# RandomNets – Implicit and Vectorized Ensemble Neural Networks



Esben Jannik Bjerrum


RDKit UGM

September 2025

Prague


# Model Ensemble – Enhancing Model Performance

- **Averaging prediction** from multiple, slightly different, models is a known trick to enhance predictivity
- Well-known models such as **Random Forest** uses ensembles of decision trees
- **Bagging**: Bootstrap Aggregation
  - Sampling the dataset with replacement gives different datasets and thus models
- **Random subspace method** (feature bagging):
  - Models are presented with different subsets of features to induce differences.
- **Efficient approach**: Tox24 blind challenge showed top contenders to be ensembles or consensus models[1]



Decision Tree

Weak model



Random Forest

Strong ensemble

# Manual or Loop Approaches

- Often ensembles approached with looping or manually training different models and later averaging
  - Efforts scale linearly with number of models
  - Can be a hassle to manage: Training, saving, loading, prediction, averaging
- .... often need to develop some sort of framework if it doesn't exist already
- But is maximum flexible....




# Vectorization for *Speed* 🏃 🏃 🏃

- Vectorization of problems is a known approach to avoid slow for-loops (especially in Python)

=> “Simply” add a new dimension to the tensors going through the neural network.

- ```
fp_resaped =  
fp.unsqueeze(2).expand(  
    -1, -1, n_ensemble_sel  
    ) # -> samples, fp_size,  
    n_nns_sel
```
- Then of course “just” adapt the network layers as well .....





# How to Create Different Models Within the Neural Network?

- “The Blind Men and the Elephant”  
🐘” TittthaSutta, Udāna 6.4, Khuddaka Nikaya  
Ca. 500 BC.  
=> Each man describe only a part of the  
elephant, by talking they discover they only  
see part of the truth  
(in other versions they start fighting over who’s  
lying 🤪)
- Similar to Random Forests we can do:
- **input masking** (Random Subspace Method)
  - Was easy to implement!
- Use only some samples to some members  
(**subsampling**, so not 1:1 to bagging)
  - Was way more challenging to implement!




# Implicit Ensembling Saves Memory (and Worked Better)

- Ran **out of memory** on my limited GPU with even low numbers of internal models (GTX 1060 with 6GB VRAM ~2017)
- Implemented **weight sharing** via Conv1D layers, turning the ensemble implicit
- Fixed **input mask**
- Output masking (**sample masking**) where each sample was only allowed backpropagation for predefined output neurons.
- Features were Morgan Fingerprints Radius 2
- Any Scikit-Mol transformer/Pipeline can be used.




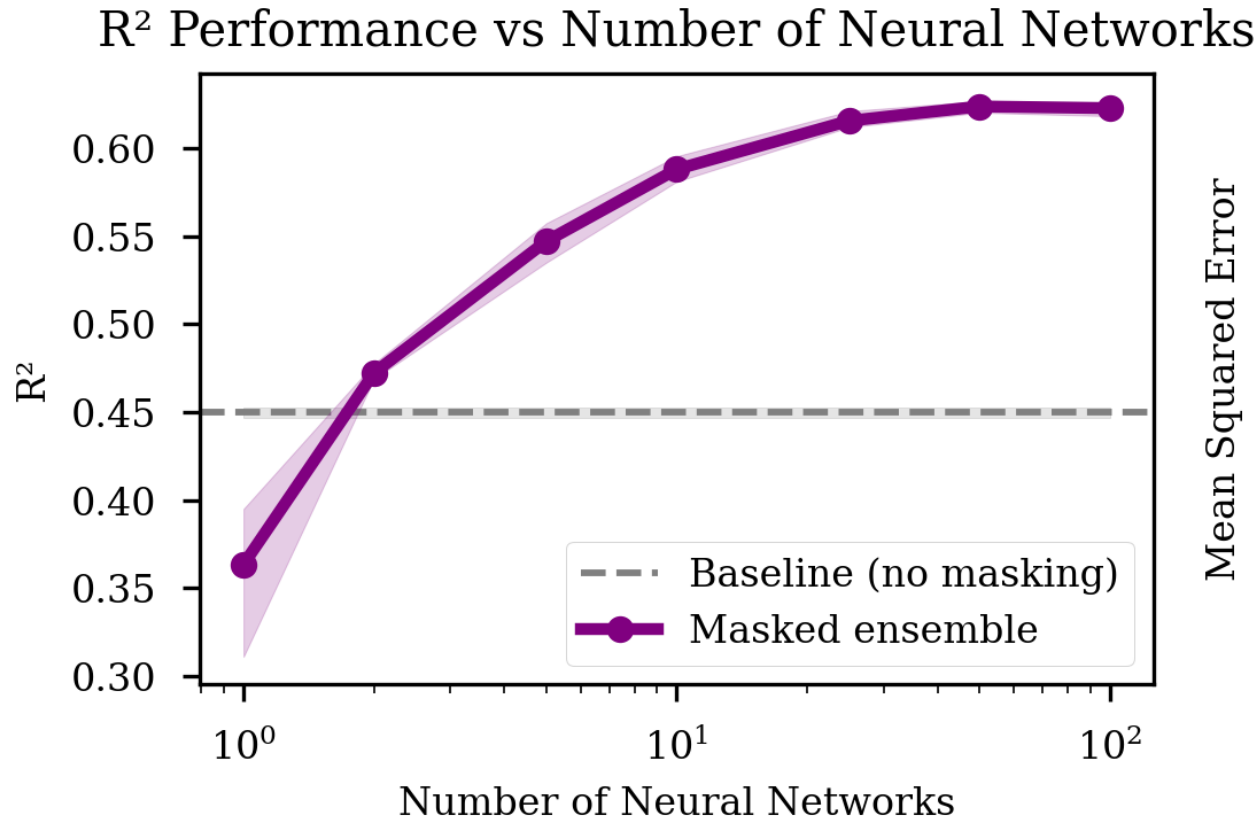
Open-Source Cheminformatics  
and Machine Learning



**I have no more excuses**



# Ensembling leads to better model predictions




Already two members in ensemble is better than single network with no masking. Effect plateaus after 50 implicit networks. Tested on ExcapeDB SLC6A3 dataset.



# Feature masking >> Sample-masking

- **Feature-masking** around 0.7 improved ensemble performance the most.
- **Sample-masking** only decreased performance
- (Inverse proportional to difficulty of implementation, 🧑🏻)

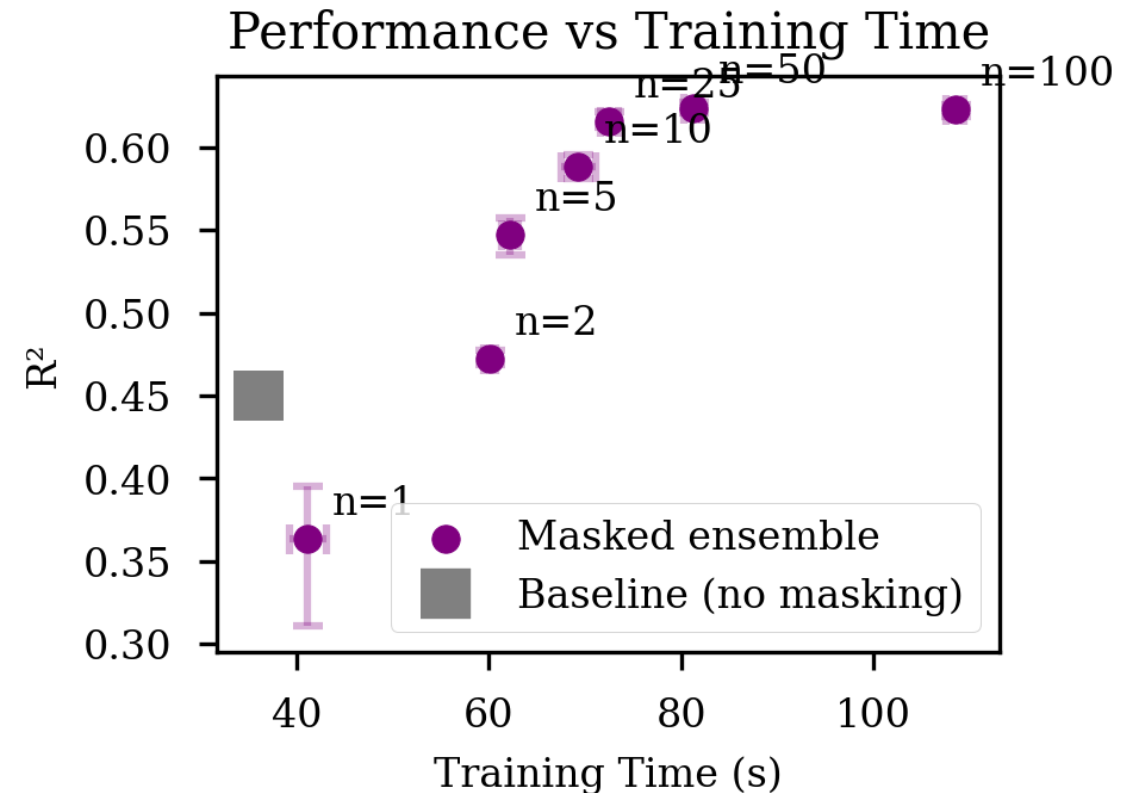
More Feature Masking



More Sample Masking

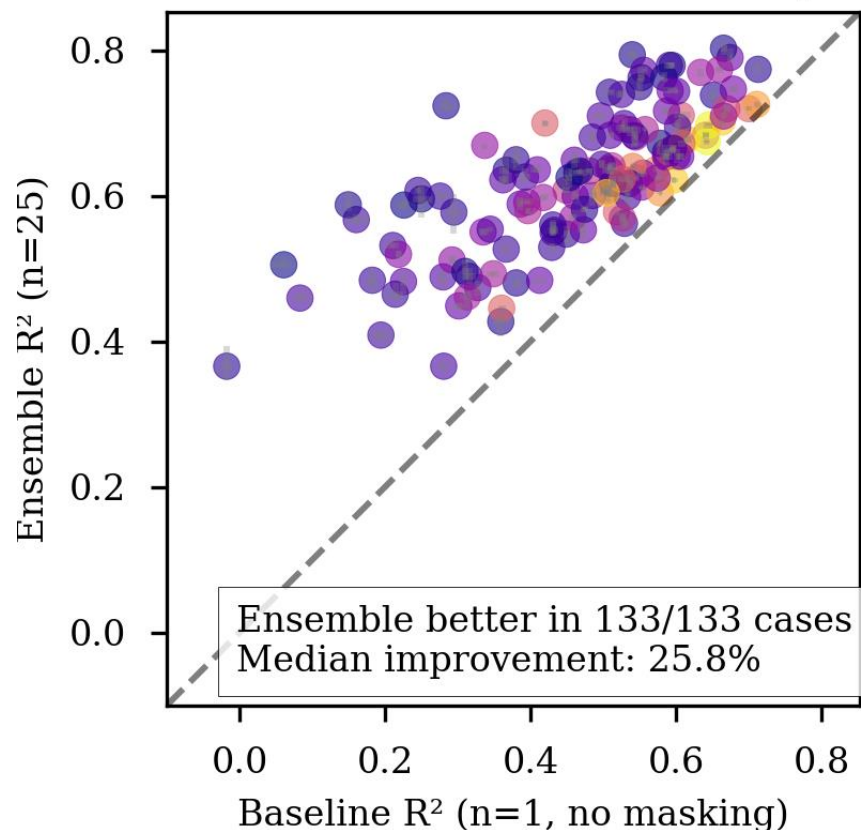
# Very Efficient Training

- Performance levels off after 50 members in ensemble.
- But corresponds to only  **$\sim 2\times$  training time** compared to single network!
- Fingerprint and  $y\_value$  tensors moved to GPU *before* expansion along  $n\_ensemble$  dimension
- Effective **minibatch size increased** significantly
- $\Rightarrow$  Maybe why we see more robust training in fewer epochs?

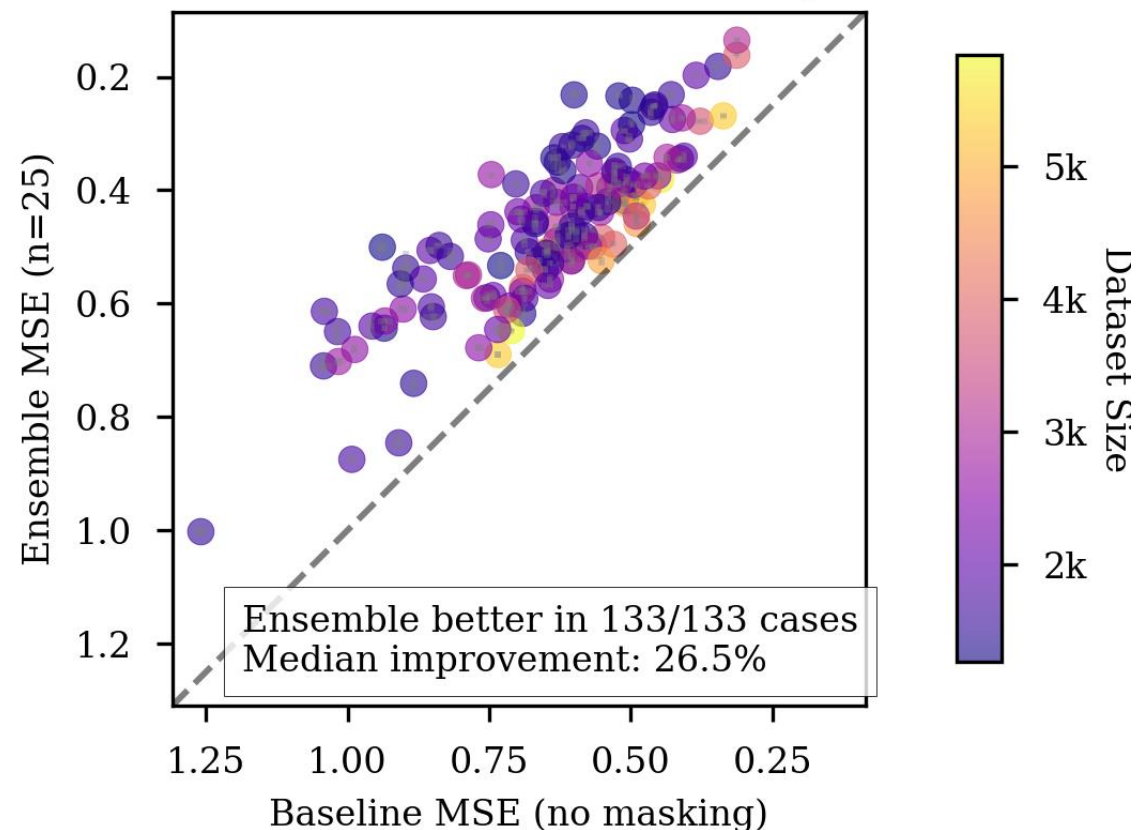


# Benchmarking on a collection of datasets

Cross-Dataset Performance Comparison




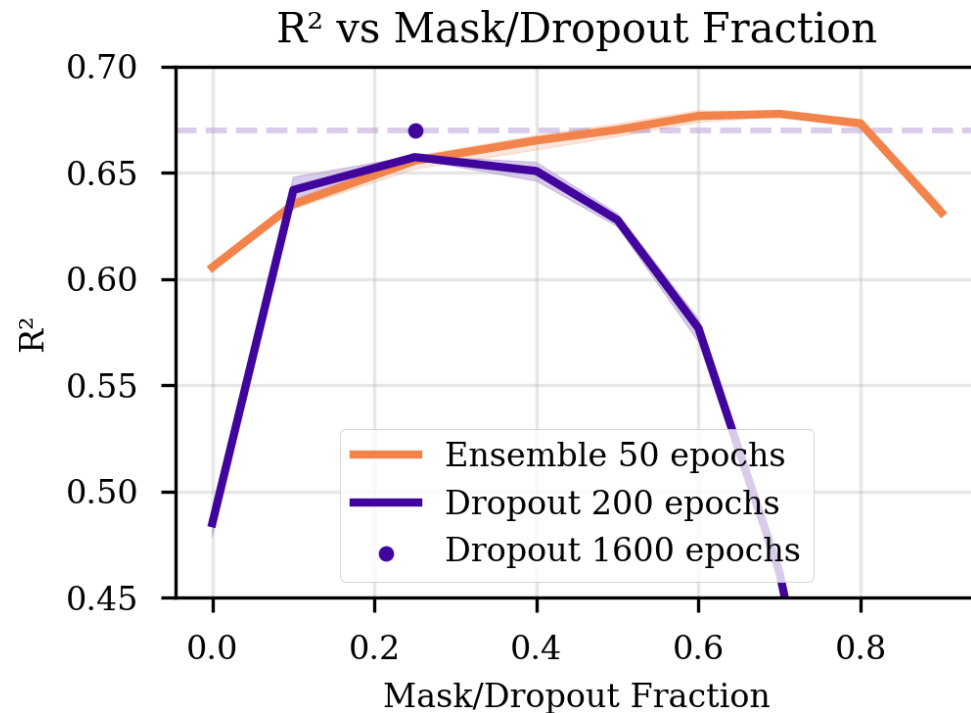
Cross-Dataset Performance Comparison



- Tests on **133** regression datasets from **ExcapeDB** (binding and activity).
- Using RandomNets ensembles **unidirectionally improved performance**, being most pronounced for smaller datasets.

Sun, Jiangming, Nina Jeliazkova, Vladimir Chupakin, et al. 'ExcapeDB: An Integrated Large Scale Dataset Facilitating Big Data Analysis in Chemogenomics'. Zenodo, 29 November 2016.

# Comparison with Dropout



- **Dropout-In** seems conceptually very similar to input-masking of implicit ensembles.
- Dropout network sampled multiple times with dropout-in active and prediction averaged.
- However, **ensemble performance could not be reached**, even with extended training to account for higher effective mini batch-size of ensemble approach.

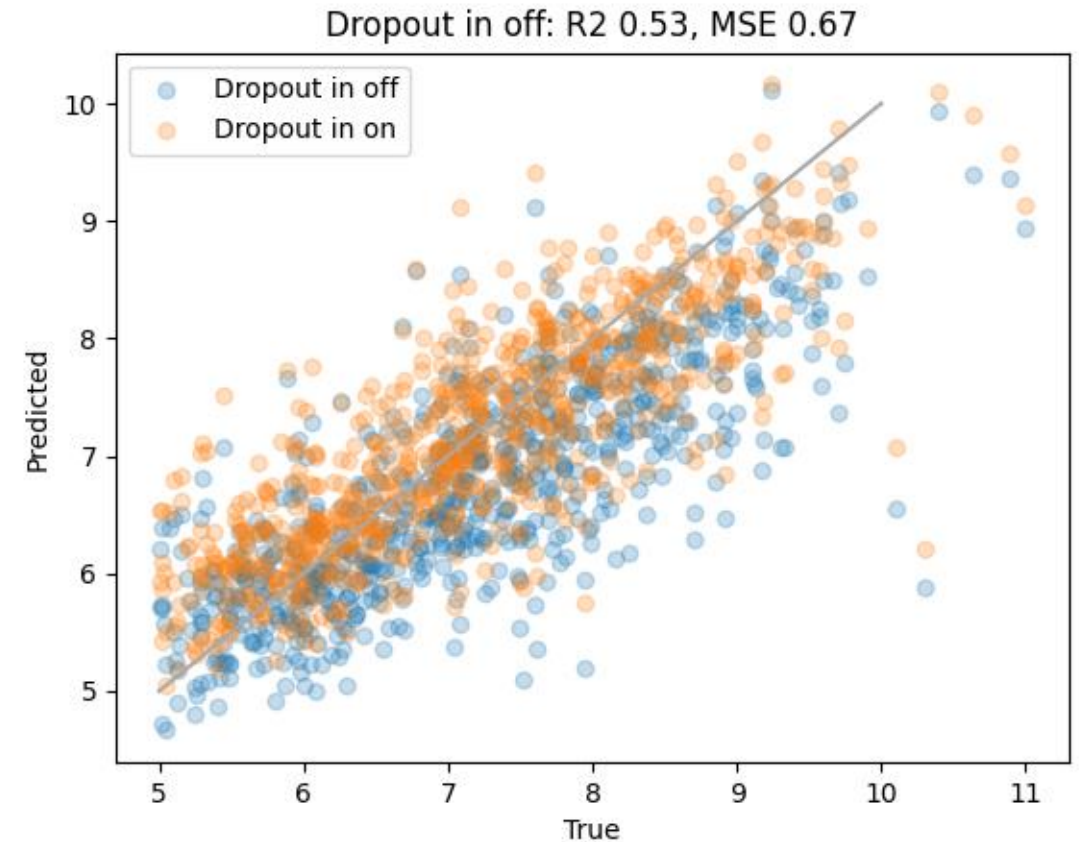


# Surprising effect of Dropout-In


- When training with dropout-in, this could **NOT** be switched off during inference without seeing a **systematic lower prediction**.

=> Dropout\_in networks was **sampled n\_ensemble times** with dropout\_in active and **prediction averaged**.

- Anyone got any idea as to why this effect comes when using dropout on input, in contrast to on layer outputs? Please reach out!



# Balancing the Mask



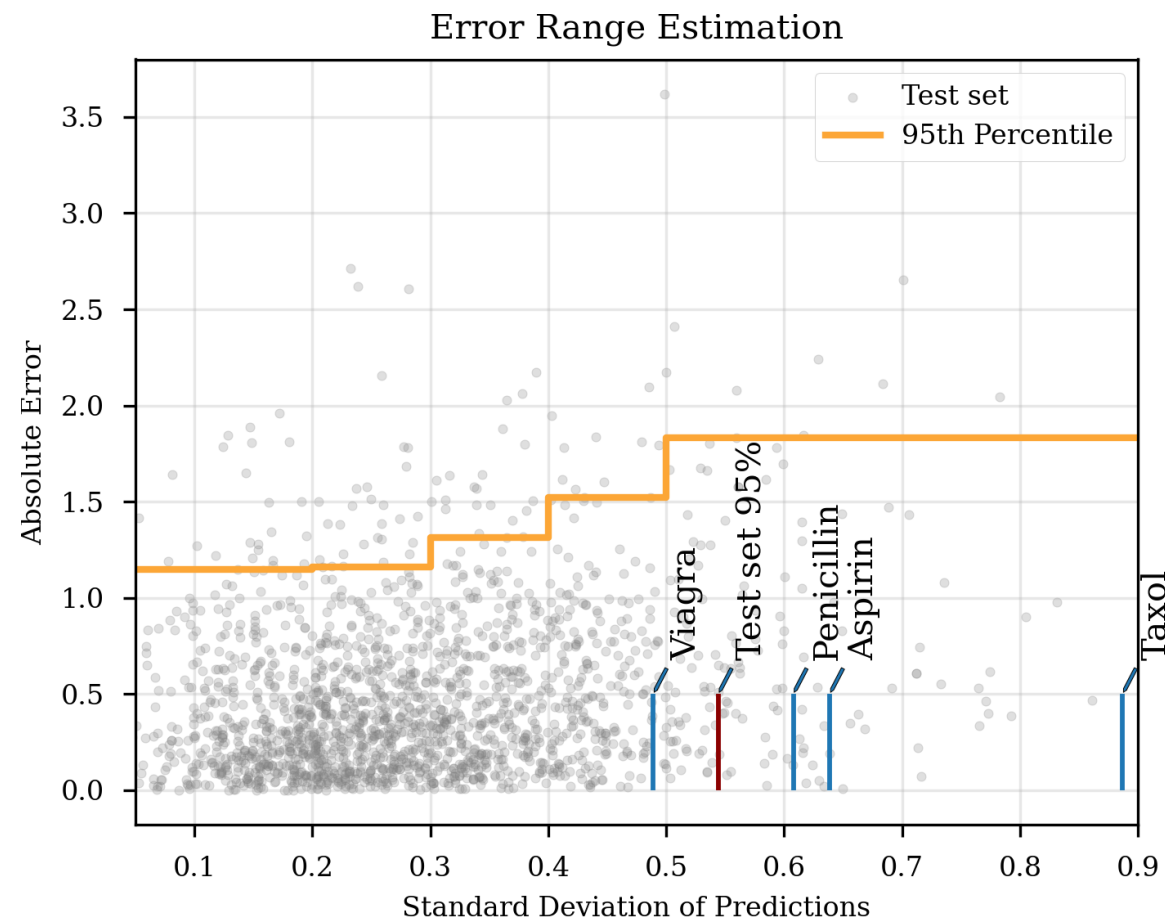
When creating the random feature mask, some features gets **more masked** than others (i.e. is used in fewer ensemble members).

Designing the input mask to have a **more even distribution** of masked features, reduces variability and improves predictive power on average.

(Results on previous slides done with fully random feature masks)


# Standard Deviation of Predictions Hints of Uncertainty and Applicability Domain

- For this dataset, SLC6A4, a serotine transporter:
- Standard deviation  $< 0.3$  have 95% of predictions below  $\sim 1.2$  units
- Standard deviations  $> 0.5$  have 95% of predictions below  $\sim 1.75$  units in absolute error
- Challenges with believed adversarial compounds associated with higher standard deviation in prediction.



# Conclusions

- **Feature masking**, but not sample masking, gave most improvement of ensemble predictions.
- Training show **sub-linear scaling**
- **Improved performance** on 133 out of 133 Bio regression datasets.
- Compared to feature dropout, shows less sensitivity to hyperparameter choice and better performance.
- RandomNets is an **attractive approach compared to single feed-forward neural networks**
- Is open-sourced at <https://github.com/EBjerrum/RandomNets>
- Currently only regression modelling but can be easily adapted (Reach out if interested).





# Acknowledgements

- Open Source Contributors of the World
- Releasers of open molecular datasets



Open-Source Cheminformatics  
and Machine Learning



# Thank you for your attention!



## Questions?